

Pherkad cheat sheet

Checks whether a draft still sounds like **you**, not a model. One voice, quoted evidence behind every score, fixes only for the sentences that drifted. Say things in your own words; the phrases below are examples, not commands.



PHERKAD

the other Guardian of the Pole
sibling of Kochab

FIRST RUN: BUILD YOUR PROFILE

`| build my voice profile`

A short interview over 3 to 5 pieces of writing you consider most you. Produces **Voice_Profile.md**: your authority, hedging, texture, sentence mechanics, structure, and tone, every marker backed by a quoted sentence of yours. Lives in your folder, never leaves it.

VOICE CHECK

`| does this sound like me | voice check this draft`

- Fingerprint: the 3 sentences most and least like you, quoted
- Seven dimensions scored 1 to 5, each with cited evidence
- The full catalog checked, each flag naming its specific tell
- Density signal: flagged constructions per 100 words (warn above 2.0), plus cluster warnings per paragraph
- Verdict: **PASS** **REVISE** **REWRITE**

THE SEVEN DIMENSIONS

- | | |
|-----------------------|---|
| 1 Grounding | where your authority comes from |
| 2 Calibration | how you hedge, how confident you get |
| 3 Texture | your concrete specifics |
| 4 Mechanics | rhythm, punctuation, fragments, humor |
| 5 AI artifacts | the shared tell catalog (right) |
| 6 Structure | how you open, build, end; what stays implicit |
| 7 Tone | the overall feel, and what you never sound like |

TARGETED REWRITES

Arrive with any REVISE or REWRITE verdict. Only the flagged sentences are rewritten, each fix explained. Facts, names, numbers, dates, sources, and emphasis stay exactly as you had them; terms of art survive.

CORRECT A FLAG

`| that one is actually me`

The correction appends to your profile with the sentence as evidence, so a later run applies it instead of repeating the flag.

TELL CATEGORIES AT A GLANCE

- 5a** filler and intensifiers ("crucial," "delve," "leverage")
- 5b** flattened stock phrases ("paradigm shift")
- 5c** antithesis family ("It doesn't just X. It Y.")
- 5d** scene-setting openers ("In a world where...")
- 5e** engagement bait ("here's the thing")
- 5f** triplet noun piling ("clarity, trust, and growth")
- 5g** counter-X constructions
- 5h** authenticity paradox ("to be honest," "frankly")
- 5i** structural artifacts: em dashes, hedge-stacking, canned closers, loaded adverbs

Quotes, technical-literal uses, and single isolated instances are exempt. Density and clustering are the tell.

CLI LINTER (NO MODEL NEEDED)

`| python3 tools/voicelint.py draft.md`

- Mechanical tells with exact line and column numbers
- `--json --strict` for CI or a pre-commit hook
- Exit 0 no error-level findings, 1 an error (or a warning under `--strict`), 2 a usage, IO, or config error
- Reads Markdown, plain text, and HTML; stdlib only, Python 3.8+

PERSONAL LINTER CONFIG

`| make me a linter config from my profile`

Your dash stance, crutch-word caps, and extra bans as a **voice_config.json** the CLI linter can run anywhere.

ABOUT THE DEFAULTS

The shipped rules are the author's own, built from tells observed across many AI-assisted documents, and they ship strict. Your register may differ: edit **tools/voice_config.json** to your purpose, or relax single rules in your own config (**tools/examples/relaxed.json** shows the shape), with the override recorded in your profile so both layers agree.

INSTALL

Claude (Cowork / claude.ai): add **pherkad.skill** via Settings > Capabilities, or build it with **./build.sh**.
Claude Code: copy **skills/pherkad/** into your skills directory.
Then ask for a voice check on any draft.

Your data: your profile is built from your writing and stays in your folder. The repo ships no personal data; its only persona is fictional.

The one rule: every verdict carries quoted evidence.
A human-written draft can fail; the test is you. Density is the tell, never one stray phrase.