Pherkad cheat sheet

Checks whether a draft still sounds like **you**, not a model. One voice, quoted evidence behind every score, fixes only for the sentences that drifted. Say things in your own words; the phrases below are examples, not commands.

PHERKAD
the other Guardian of the Pole
sibling of Kochab

FIRST RUN: BUILD YOUR PROFILE

`build my voice profile`

A short interview over at least four pieces of writing you consider most you. Produces **Voice_Profile.md**: your authority, hedging, texture, sentence mechanics, structure, and tone, every marker backed by a quoted sentence of yours. Lives in your folder, never leaves it.

VOICE CHECK

`does this read as me` `audit this`

- **Quick** (the default): one table, each flagged passage ruled, then a verdict line
- **Full audit**: the 3 sentences most and least like you, seven dimensions scored with evidence, the full catalog
- Density above 2.0 flags per 100 words (150 words and 3 flags or more), plus cluster warnings
- Verdict: PASS REVISE REWRITE

THE SEVEN DIMENSIONS

1 **Grounding**	where your authority comes from
2 **Calibration**	how you hedge, how confident you get
3 **Texture**	your concrete specifics
4 **Mechanics**	rhythm, punctuation, fragments, humor
5 **AI artifacts**	the shared tell catalog (right)
6 **Structure**	how you open, build, end; what stays implicit
7 **Tone**	the overall feel, and what you never sound like

TARGETED REWRITES

Arrive with any REVISE or REWRITE verdict. Only the flagged sentences are rewritten, each fix explained. Facts, names, numbers, dates, sources, emphasis, and terms of art must stay as you had them; a rewrite that changes them is wrong, so check that it did not.

CORRECT A FLAG

`that one is actually me`

The correction appends to your profile with the sentence as evidence, so a later run applies it instead of repeating the flag.

TELL CATEGORIES AT A GLANCE

5a	filler and intensifiers ("crucial," "delve," "leverage")
5b	flattened stock phrases ("paradigm shift")
5c	antithesis family ("It doesn't just X. It Y.")
5d	scene-setting openers ("In a world where...")
5e	engagement bait ("here's the thing")
5f	triplet noun piling ("clarity, trust, and growth")
5g	counter-X constructions
5h	authenticity paradox ("to be honest," "frankly")
5i	structural artifacts: em dashes, hedge-stacking, canned closers, loaded adverbs

Quotes, technical-literal uses, and single isolated instances are exempt. Density and clustering are the tell.

CLI LINTER (NO MODEL NEEDED)

`python3 tools/pherkad.py check --surface email draft.md`

- Both mechanical engines, one list, a stable rule id on every finding
- **--format json** or **sarif**, **--strict** for CI; **--decisions FILE** keeps a ruled finding quiet
- Exit 0 no error-level findings, 1 an error (or a warning under **--strict**), 2 a usage, IO, or config error
- Reads Markdown, plain text, and HTML; stdlib only, Python 3.8+

PERSONAL LINTER CONFIG

`make me a linter config from my profile`

Your dash stance, crutch-word caps, and extra bans as a **voice_config.json** the CLI linter can run anywhere.

ABOUT THE DEFAULTS

The shipped rules are the author's own, built from tells observed across many AI-assisted documents, and they ship strict. Your register may differ: edit **tools/voice_config.json** to your purpose, or relax single rules in your own config (**tools/examples/relaxed.json** shows the shape), with the override recorded in your profile so both layers agree.

INSTALL

Claude (Cowork / claude.ai): add **pherkad.skill** via Settings > Capabilities, or build it with **./build.sh**.
Claude Code: copy **skills/pherkad/** into your skills directory.
Then ask for a voice check on any draft.

Your data: your profile is built from your writing and stays in your folder. The repo ships no personal data; its only persona is fictional.

The one rule: every verdict carries quoted evidence. A human-written draft can fail; the test is you. Density is what gives it away, never one stray phrase.